

Name:Vikas Mane
Class:BE-B
Roll Number:B1921152
PRN Number:72000291F

Experiment No: 6

Title: Vector and Matrix Operations-

Design parallel algorithm to

1. Add two large vectors
2. Multiply Vector and Matrix
3. Multiply two $N \times N$ arrays using n_2 processors

Aim: Implement $n \times n$ matrix parallel addition, multiplication using CUDA, use shared memory.

Prerequisites:

- Concept of matrix addition, multiplication.
- Basics of CUDA programming

Objectives:

Student should be able to learn parallel programming, CUDA architecture and CUDA processing flow

Theory:

A straightforward matrix multiplication example that illustrates the basic features of memory and thread management in CUDA programs

- Leave shared memory usage until later
- Local, register usage
- Thread ID usage
- $P = M * N$ of size $WIDTH \times WIDTH$
- One thread handles one element of P
- M and N are loaded $WIDTH$ times from global memory

Matrix Multiplication steps

1. Matrix Data Transfers
2. Simple Host Code in C
3. Host-side Main Program Code
4. Device-side Kernel Function
5. Some Loose Ends

Step 1: Matrix Data Transfers

```
// Allocate the device memory where we will copy M
to Matrix Md;
Md.width = WIDTH;
Md.height = WIDTH;
Md.pitch = WIDTH;
int size = WIDTH * WIDTH * sizeof(float);
cudaMalloc((void**)&Md.elements, size);
// Copy M from the host to the device
cudaMemcpy(Md.elements, M.elements, size, cudaMemcpyHostToDevice);
// Read M from the device to the host into P cudaMemcpy(P.elements,
Md.elements, size, cudaMemcpyDeviceToHost);
...
// Free device memory
cudaFree(Md.elements);
```

Step 2: Simple Host Code in C

```
// Matrix multiplication on the (CPU) host in double precision
```

```
// for simplicity, we will assume that all dimensions are equal
```

```
void MatrixMulOnHost(const Matrix M, const Matrix N, Matrix P)
{
    for (int i = 0; i < M.height; ++i)
        for (int j = 0; j < N.width; ++j) { double sum = 0;
            for (int k = 0; k < M.width; ++k) {
                double a = M.elements[i * M.width + k]; double b = N.elements[k * N.width +
                j]; sum += a * b;
            }
            P.elements[i * N.width + j] = sum;
        }
}
```

Multiply Using One Thread Block

- One Block of threads compute matrix P – Each thread computes one element of P
- Each thread – Loads a row of matrix M
 - Loads a column of matrix N
 - Perform one multiply and addition for each pair of M and N elements
 - Compute to off-chip memory access ratio close to 1:1 (not very high)
- Size of matrix limited by the number of threads allowed in a thread block

Step 3: Host-side Main Program Code

```
int main(void) {
    // Allocate and initialize the matrices
    Matrix M = AllocateMatrix(WIDTH, WIDTH, 1); Matrix N = AllocateMatrix(WIDTH,
    WIDTH, 1); Matrix P = AllocateMatrix(WIDTH, WIDTH, 0);
    // M * N on the device MatrixMulOnDevice(M, N, P);
    // Free matrices FreeMatrix(M)
    ;
    FreeMatrix(N)
    ;
}
```

```
FreeMatrix(P);

return 0;

}
```

Host-side code

```
// Matrix multiplication on the device
void MatrixMulOnDevice(const Matrix M, const Matrix N, Matrix P)
{
    // Load M and N to the device Matrix Md = AllocateDeviceMatrix(M);
    CopyToDeviceMatrix(Md, M);
    Matrix Nd = AllocateDeviceMatrix(N); CopyToDeviceMatrix(Nd, N);
    // Allocate P on the device
    // Setup the execution configuration dim3 dimBlock(WIDTH, WIDTH);
    dim3 dimGrid(1, 1);
    // Launch the device computation threads! MatrixMulKernel<<<dimGrid,
    dimBlock>>>(Md, Nd, Pd);
    // Read P from the device CopyFromDeviceMatrix(P, Pd);
    // Free device matrices

    FreeDeviceMatrix(Md); FreeDeviceMatrix(Nd); FreeDeviceMatrix(Pd);
}
```

Step 4: Device-side Kernel Function

```
// Matrix multiplication kernel - thread specification
global void MatrixMulKernel(Matrix M, Matrix N, Matrix P)
{
    // 2D Thread ID
    int tx = threadIdx.x; int ty = threadIdx.y;
    // Pvalue is used to store the element of the matrix
    // that is computed by the thread float Pvalue = 0;
    for (int k = 0; k < M.width; ++k)
    {
        float Melement = M.elements[ty * M.pitch + k]; float Nelement =
        Nd.elements[k * N.pitch + tx]; Pvalue += Melement * Nelement;
    }
    // Write the matrix to device memory;
    // each thread writes one element P.elements[ty * P.pitch + tx] = Pvalue;
}
```

Step 5: Some Loose Ends

- Free allocated CUDA memory

Facilities:

Latest version of 64 Bit Operating Systems, CUDA enabled NVIDIA Graphics card

Input:

Two matrices

Output:

```

uniform_int_distribution<int> distribution(0, 1000);
for (int i = 0; i < M; ++i) {
    for (int j = 0; j < N; ++j) {
        a[i][j] = distribution(generator);
    }
}

for (int i = 0; i < N; ++i) {
    for (int j = 0; j < P; ++j) {
        b[i][j] = distribution(generator);
    }
}

double t1, t2;

printf("Time for matrix multiplication of two %dx%d matrices\n", M, N);
for (int num_threads = 1; num_threads <= 10; ++num_threads)
// set the number of threads to num_threads
omp_set_num_threads(num_threads);
memset(c, sizeof(c), 0);
// time before starting matrix multiplication
t1 = omp_get_wtime();
#pragma omp parallel reduction(+: c)
{
    #pragma omp for collapse(3)
    for (int i = 0; i < M; ++i) {
        for (int j = 0; j < P; ++j) {
            for (int k = 0; k < N; ++k) {
                c[i][j] += (a[i][k] * b[k][j]);
            }
        }
    }
}
// time after finishing matrix multiplication
t2 = omp_get_wtime();
printf("Using %d thread(s): %g\n", num_threads, t2 - t1);

```

```

sahilbansal@ubuntu:~/Desktop/OpenMP$ g++ matrixMul
matrixMul
matrixMultiplication
matrixMultiplication.cpp
sahilbansal@ubuntu:~/Desktop/OpenMP$ g++ matrixMul
tiplication.cpp -o matrixMul -fopenmp
sahilbansal@ubuntu:~/Desktop/OpenMP$ ./matrixMul
Time for matrix multiplication of two 1000x1000 ma
trices using different no. of threads:
Using 1 thread(s): 8.43642
Using 2 thread(s): 4.09487
Using 3 thread(s): 3.69263
Using 4 thread(s): 3.76583
Using 5 thread(s): 4.36293
Using 6 thread(s): 4.43307
Using 7 thread(s): 3.95484
Using 8 thread(s): 4.66723
Using 9 thread(s): 4.91912
Using 10 thread(s): 4.19519
sahilbansal@ubuntu:~/Desktop/OpenMP$

```

Conclusion:

We learned parallel programming with the help of CUDA architecture.