

Title: Mushroom classification

Code:

```
import numpy as np # Linear algebra
import pandas as pd
import os
for dirname, _, filenames in os.walk('/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
mushroom = pd.read_csv('mushrooms.csv')
mushroom.info()
mushroom.head(5).transpose()
mushroom.isna().sum()
mushroom.duplicated().sum()
for i in mushroom.columns:
    print(mushroom[str(i)].value_counts())
    print('\n')
mushroom_en = pd.get_dummies(mushroom, drop_first = True)
mushroom_en.shape
mushroom_en.columns
X_feature = list(mushroom_en.columns)
X_feature.remove('class_p')
X = mushroom_en[X_feature]
X.columns
Y = mushroom_en['class_p']
Y.head()
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,Y, test_size =
0.30, random_state = 42)
print("X_train.shape",X_train.shape)
print("X_test.shape",X_test.shape)
print("y_train.shape",y_train.shape)
print("y_test.shape",y_test.shape)
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors = 21)
knn.fit(X_train, y_train)

#predicting for test data
knn_test_pred = knn.predict(X_test)
#evaluate the model
from sklearn.metrics import accuracy_score
print(accuracy_score(y_test, knn_test_pred))
pd.crosstab(y_test, knn_test_pred, rownames = ['Actual'], colnames=
['Predictions'])
#error on train data
knn_train_pred = knn.predict(X_train)
print(accuracy_score(y_train, knn_train_pred))
```

```

pd.crosstab(y_train, knn_train_pred, rownames = ['Actual'], colnames =
['Predictions'])
acc = []

# running KNN algorithm for 3 to 50 nearest neighbours(odd numbers) and
# storing the accuracy values
for i in range(3,50,2):
    neigh = KNeighborsClassifier(n_neighbors=i)
    neigh.fit(X_train, y_train)
    train_acc = np.mean(neigh.predict(X_train) == y_train)
    test_acc = np.mean(neigh.predict(X_test) == y_test)
    acc.append([train_acc, test_acc])
import matplotlib.pyplot as plt # library to do visualizations
# train accuracy plot
plt.plot(np.arange(3,50,2),[i[0] for i in acc],"ro-")

```

Output:

1. Mushroom data analysis

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8124 entries, 0 to 8123
Data columns (total 23 columns):
#   Column              Non-Null Count  Dtype
---  -
0   class                8124 non-null  object
1   cap-shape            8124 non-null  object
2   cap-surface          8124 non-null  object
3   cap-color            8124 non-null  object
4   bruises              8124 non-null  object
5   odor                8124 non-null  object
6   gill-attachment      8124 non-null  object
7   gill-spacing         8124 non-null  object
8   gill-size            8124 non-null  object
9   gill-color           8124 non-null  object
10  stalk-shape          8124 non-null  object
11  stalk-root           8124 non-null  object
12  stalk-surface-above-ring 8124 non-null  object
13  stalk-surface-below-ring 8124 non-null  object
..

```

Out[15]:

Predictions	0	1
Actual		
0	1257	0
1	3	1178

Out[17]:

Predictions	0	1
Actual		
0	2951	0
1	5	2730

2. First five rows of dataset

Out[5]:

	0	1	2	3	4
class	p	e	e	p	e
cap-shape	x	x	b	x	x
cap-surface	s	s	s	y	s
cap-color	n	y	w	w	g
bruises	t	t	t	t	f
odor	p	a	i	p	n
gill-attachment	f	f	f	f	f
gill-spacing	c	c	c	c	w
gill-size	n	b	b	n	b
gill-color	k	k	n	n	k
stalk-shape	e	e	e	e	t
stalk-root	e	c	c	e	e
stalk-surface-above-ring	s	s	s	s	s
stalk-surface-below-ring	s	s	s	s	s
stalk-color-above-ring	w	w	w	w	w
stalk-color-below-ring	w	w	w	w	w
veil-type	p	p	p	p	p
veil-color	w	w	w	w	w
ring-number	o	o	o	o	o
ring-type	p	p	p	p	e
spore-print-color	k	n	n	k	n
population	s	n	n	s	a

3. Checking for null values

```
Out[6]: class          0
        cap-shape      0
        cap-surface    0
        cap-color      0
        bruises        0
        odor           0
        gill-attachment 0
        gill-spacing   0
        gill-size      0
        gill-color     0
        stalk-shape    0
        stalk-root     0
        stalk-surface-above-ring 0
        stalk-surface-below-ring 0
        stalk-color-above-ring 0
        stalk-color-below-ring 0
        veil-type      0
        veil-color     0
        ring-number    0
        ring-type      0
        spore-print-color 0
        population     0
        habitat        0
        dtype: int64
```

4. Accuracy graph

```
Out[18]: [matplotlib.lines.Line2D at 0x29d8fa3a7f0]
```

